

Company News Shocks in Polymarket

Late-September data review

Avi Kabra · Applied Mathematics Senior Thesis · Yale University

Do prediction-market prices re-price when company news lands, and can we
measure it cleanly?

Since August, the pivot is done

- **Was:** sports, politics, and geopolitics; headline-only text; a preliminary null result.
- **Now:** one area, individual-company contracts, with full real trade history and a complete log-odds analysis.
- Built the universe, pulled every trade, audited the news, and tested matching directly.

Q1: Usable sample of 6,928 contracts across 86 companies



- Individual companies only (public, plus a few private ladders). No sports, politics, crypto, or index markets.
- Window: April 2024 to August 2026.
- Two contract types: monthly price **strike ladders** and one-off **corporate events**.

Contract family	Contracts
Price ladders (monthly strikes)	6,174
Corporate events (earnings, M&A, executives)	470
Revenue ladders	170
Valuation ladders (private companies)	110
Market-cap ladders	4

Q1: A liquid core, but concentrated

Liquidity floor	Share of contracts
Traded at all	93.6%
Over \$1,000 volume	70.4%
Over \$10,000 volume	21.0%
Top 500 markets	62.6% of all volume

- Total lifetime volume: **\$79M**.
- **Clean set** (at least 200 trades and 10 active trading days): **473 contracts**.





Q2: A real reaction is mostly nothing, then a big jump

- Measured on **72.5M** minute-by-minute observations from the liquid contracts.

Six-hour log-odds change	Value
Windows that are zero-move	74.6%
Distribution mean	about 0
Distribution std. dev.	0.44

- Most of the time the market does not trade, so it does not move. This is the central fact.




Q2: When it moves, it moves a lot

Non-zero six-hour move	Log-odds
Median	0.41
90th percentile	1.66
99th percentile	5.17

- Reference: 0.40 to 0.44 is +0.16 log-odds. A 0.41 move is a decisive probability swing.
- The reaction is real and economically large. It is just sparse.

Q2: Demeaning and normalizing cleans the series

- Removing each contract's mean and scaling to unit variance removes level differences across contracts (different strikes and regimes).
- Every contract is then on one comparable scale, expressed in standard-deviation moves.
- The standardized series is the form the model consumes.




Q3: Assessing GDELT as the article source

- The free option (GDELT), tested directly on 40 real corporate events:

Test	Result	Verdict
Event coverage (within 2 days)	78%	partial
Primary journalism share	1.2%	buried in noise
Body-text scrape yield	69%	usable, biased
Market resolution vs. timestamp	24 min vs. 1 day	about 60x too coarse

GDELT validation probe — is GDELT+scraping enough as the news corpus?



Q3: Matching approach

WHAT'S THERE

- Naive name-matching is about 5% correct on a major company.
- GDELT supplies candidate links; usable article text recovers about 69% of the time.

NEXT STEPS

- Filter candidates by company, ticker, time, and contract dates.
- A small model judges relevance; validate on a hand-audited sample.
- Add a licensed news feed for the precision pass.

What changed since the preliminary run

	Preliminary (August)	This review
Universe	sports, politics, geopolitics	one company-equity universe
Prices	partial, headline-timed	full real trade history
Zeros	66%	74.6%

Next Steps

- Build the matching pipeline on the free corpus; hand-audit a sample of matches.
- Define the reaction target and rerun on the clean set.
- Run LSTM, Transformer, TCN, and cross-category comparisons.